

Implementation of a Brute Force Attack on A5/1 Cryptographical Algorithm in a GPU-based Volunteering Computing Project


Bulavintsev V., Semenov A., and Zaikin O.

Matrosov Institute for System Dynamics and Control Theory of Siberian Branch of Russian Academy of Sciences (ISDCT SB RAS)

Abstract. We present an advanced brute force attack on the A5/1 cryptographical algorithm, that is still broadly used in modern GSM networks. We use a well-known idea introduced by R.Anderson more than 20 years ago to greatly reduce the search space. The primary contribution of this article is the implementation of the Anderson's attack on a GPU platform with the bitslice technique. The preliminary estimates of the attack's speed showed that, with the use of a GPU's processing power, the attack could be conducted in the real time on a modern computer cluster or in a volunteer computing project. To verify our estimates with the use of BOINC technology we conceived the specialized volunteering computing project and executed our variant of Anderson's attack within it. In the 7 days the project lasted, 10 A5/1 cryptanalysis problems were solved. The results presented in this work provide yet another proof of A5/1's cryptographical weakness that make it totally unsuitable for transmission of any kind of sensitive data through modern GSM networks.

Keywords: A5/1 algorithm, brute-force cryptanalysis, GPU, volunteering computing

1 Introduction

A5/1 algorithm is a keystream generator with key length of 64 bits. It is used to encrypt voice and SMS traffic in 2nd generation ("2G") GSM networks. 3rd generation GSM networks accept 2G communication protocol for backward compatibility. A practice of sending voice traffic through 2G protocols in 3G networks to conserve bandwidth  increase service availability is widely adopted among mobile phone operators. A widely adopted practice among mobile operators is to still use 2G protocols to route voice traffic even in 3rd generation GSM networks.

One can name A5/1 as one of the most publicly recognized cryptographical algorithms, along with RSA and DES, it's discussion reaching far beyond the borders of professional cryptographers community. For example, article [?] publicly examines NSA's ability to efficiently decrypt A5/1. So, in this article we won't touch on the details of A5/1's creation history and the reasons of it's rise as the

world's 'de facto' mobile communication standard. Milestones in cryptanalysis of A5/1 would be briefly outlined in Related Works section.

Among all the different methods of A5/1 cryptanalysis we distinguish those which were realized in practice and allowed to reliably conduct the cryptanalysis procedure on an non-weakened variant of algorithm. Apparently, the first such attack was performed in 2008 with the help of the special FPGA-based computational platform COPACOBANA [?]. In 2009 distributed algorithms for boolean satisfiability (SAT) were used to solve several cryptanalysis instances of stock A5/1 in a specialized grid system 'BNB-Grid' [?]. These results were further improved in 2011 [?]. By the end of 2009 the 'A5/1 Cracking Project' had published rainbow tables [?] for A5/1. Provided with 8 bursts (912 bits) of keystream these tables allowed to find the secret key in less than a minute with more than 85% probability. Despite the huge (over 2 Tb) size, to this day these rainbow tables provide the most practical method of A5/1 cryptanalysis. It's main shortcoming is that the probability of success is significantly less than 100%. Meanwhile, the growth of computational power of GPUs and FPGAs made practical the attack based on reduction of search space from 2^{64} to 2^{53} , which was described by R. Anderson in 1994 [?]. As was mentioned earlier, FPGA-based variant of the Anderson's attack was already performed in 2008 by COPACOBANA creators. So, the primary goal of our work is to demonstrate the viability of GPU-based variant of the attack. Let us note that GPUs are much easier to operate than FPGAs, and the former belong to the class of consumer-grade devices and could be found in any modern PC, while the latter belong to the class of specialized equipment. With the usage of the BOINC software platform [?], this qualities of GPUs allowed us to implement the attack in the form of a volunteering computing project using idle computational capabilities of the project members home PCs. Our estimates of attack's speed were based on our previous work [?].

Let's make a brief outline of the article's contents. In Section 2 we describe the A5/1 algorithm along with some advanced brute-force attacks on it. Section 3 introduces bit-slicing technique and goes through important details of implementing the Anderson's attack with it. Section 4 provides a look into the internal organization of the volunteering project we created for conducting the attack with GPUs. ?? contains the retrospective of A5/1 cryptanalysis works related to our study.

2 A5/1 algorithm and some attack on it

A5/1 keystream generator consists of 3 linear feedback shift registers (LFSRs) [?], defined by the following primitive polynomials:

$$\begin{aligned} X^{19} + X^{18} + X^{17} + X^{14} + 1, LFSR1; \\ X^{22} + X^{28} + 1, LFSR2; \\ X^{23} + X^{22} + X^{21} + X^8 + 1, LFSR3. \end{aligned}$$

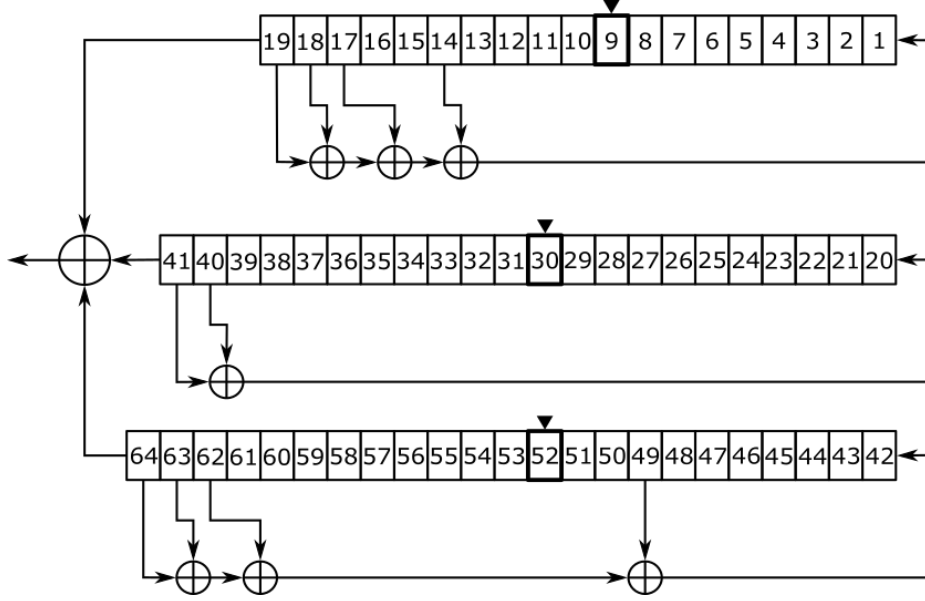


Fig. 1. A5/1 generator work scheme

The illustration of A5/1 generator's work scheme can be seen at ??.

The outputs of LFSRs are mixed with linear function, which provides the perfect correlation immunity. Non-linearity of the equations is achieved by clocking the registers asynchronously - for each clocking of the whole generator any of the 3 registers could be clocked or it could retain it's current state. Register with index $j \in \{1, 2, 3\}$ will be clocked if the following boolean function χ_j becomes 1:

$$\chi_j = (b_j \equiv \text{majority}(b_1, b_2, b_3));$$

$$\text{majority}(A, B, C) = (A \wedge B) \vee (A \wedge C) \vee (B \vee C).$$

Here b_1, b_2, b_3 denote clocking bits marked at ?? by black wedges. Conversely, if at some moment $\chi_j = 0$, LFSR j won't clock (it will remain in it's last state).

Generator A5/1 is used in the GSM protocol for high-speed encryption of large volumes of information with the short secret key. The whole process is split into 'sessions' about 3,5 hours long. Each session uses it's own 'session key'. We won't touch on the topic of specialized protocols used in GSM networks to build and transfer the session key.

Along the message data, GSM protocol sends error correction data for the message. The message complete with error correction data constitutes a 456 bits long 'frame', which is further broken down into 4 'bursts' of 114 bits. The bursts are then encrypted and send over the air. To encrypt a burst the A5/1 generator is initialized with the 64-bit 'local key', that is built using session key and a natural number called 'the frame number' (FN). After the encryption of

one burst is complete, the frame number is increased by 1. When FN overflows, the session ends and a new session is initialized (hence the 3,5 session length). FN is always known from the open data transmitted over the network.

If the message length is less than 23 bytes, the message would be filled with fixed pattern padding to 23 bytes. Some GSM technical messages used during the voice call are fixed length and always padded. As padding is always the same, and it is encrypted by some local key, we got a typical plain text attack scenario at hand [?]. Indeed, the padding plays the role of the known plaintext, allowing the attacker to get the corresponding keystream fragment. This vulnerability allows the attacker to get no less than two frames (912 bits) of the keystream. It was demonstrated in [?] that the knowledge of even one local key and FN is enough to efficiently restore the session key, which makes the decryption of the whole session possible.

The described vulnerability of GSM protocol allows to build some successful attacks on it, based on the idea of 'advanced brute force'. In particular, the large-scale preprocessing made possible the creation of the rainbow tables, which, provided 8 bursts of keystream, allows to determine the session key with probability in less than a minute with $> 85\%$ probability. The tables take about 2Tb of disk space. This attack is presented in detail in [?], and still stands among the most practical ones. Instead, we will focus on the idea of an attack that was suggested by Ross Anderson in 1994 in a small essay on the A5/1 cryptographic strength [?]. Next we describe the essence of the Anderson's attack.

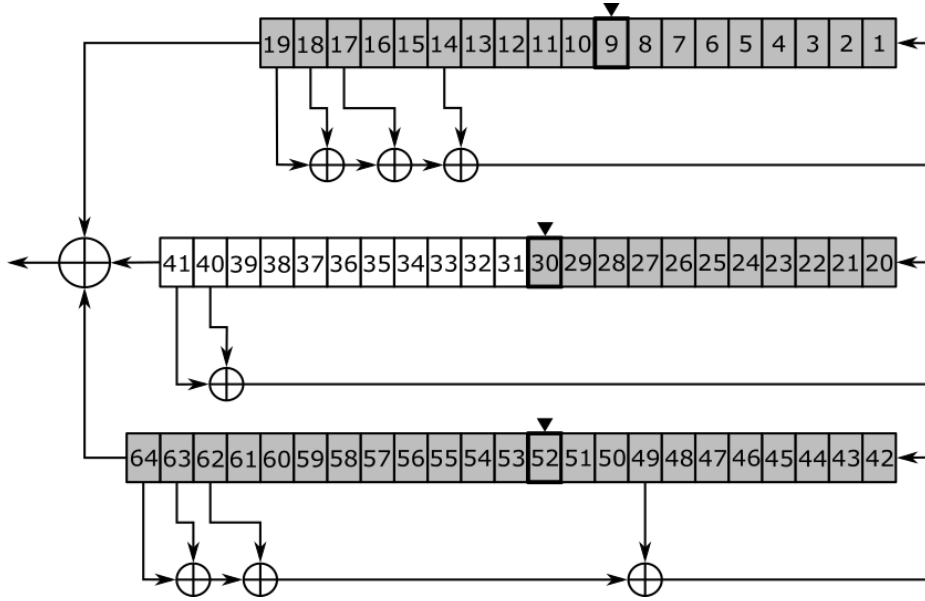


Fig. 2. The set of guessing bits used in Anderson's attack (greyed out).

Anderson's attack is a typical example of a 'guess and determine attack' (see [?]). Suppose we know the bits filling the 1st and the 3rd LFSR, and bits of the 2nd LFSR from the beginning of the register to the clocking bit (bits 31 to 41, see Fig. 2). Next suppose we know 64 bits of the keystream. Now, as was shown by R. Anderson, 11 unknown bits of the 2nd LFSR could be figured out without any additional guesses. This happens because the clocking bits are known (and so is the clocking schedule for the next 11 clockings of the 2nd LFSR), and known are 2 out of 3 XOR-ed LFSR output bits and the result of the XOR operation (from the keystream). Therefore, one can efficiently work out the unknown bits of the LFSR 2 one by one, by clocking the generator and applying XOR operation to corresponding keystream bits and output bits of LFSRs 1 and 3.

The description of the algorithm of determination of unknown 11 bits of LFSR 2 provided above makes obvious the possibility to mount a brute force attack on the A5/1 generator over the search space of 2^{53} . The simplicity of the algorithm provides an opportunity to implement it on a specialized computational architecture. One such implementation was built with FPGAs by authors of [?]. In the following sections we describe our implementation of this attack for a modern GPU.

3 Implementation of the Anderson's Attack in bit-slicing technique

The efficiency of a brute force attack is defined by two parameters: the speed of checking of the key candidates and the size of the search space. R. Anderson's idea described above give us a

$$2^{53}$$

search space for the A5/1 generator. Instead of using the 'naive' implementation of A5/1 generator, to speed up the key candidates checking procedure one can opt to use more sophisticated alternatives. We evaluated the performance of two different fast implementations of A5/1 generator in [?]. The first one was based on an idea of precomputation of the states of the registers LFSR 1-3, and keeping the states in RAM, to use them as the look-up tables. This approach demonstrated a considerable speed-up against 'naive' implementation. A similar method of precomputation of LFSRs was described in [?]. However, in [?] we found its performance inferior to the second implementation, that is based on the 'bit-slicing technique'. Next we briefly describe the idea of this technique, and its application to Anderson's attack.

Modern general-purpose computational architectures are suboptimal for implementation of many algorithms. The last operate with bits, while CPUs and GPUs typically operate with 32-256 bit words. As the result of this discrepancy, the 'naive' implementation of a cryptographical algorithm won't be fully using the platform's computational resources. For example, to conduct a 'addition modulo 2' (XOR) operation on 2 boolean arguments, a modern CPU will position the arguments values into the lowest bits of two 32-bit general-purpose

register (GPR), perform the computation and write the result into the third 32-bit GPR. In effect, 31 out of 32 bits of GPRs stay idle during this operation. This problem could be solved by using bitwise operations that work on GPRs like on boolean vectors - digit by digit. A cryptographical algorithm could be oftentimes represented as a logical scheme built with the basic logical gates. Coupled with bitwise operations, this representation allows the computational platform to apply the algorithm to as many sets of data, as there are bits in it's GPR (e.g. the platform's register capacity). This method is called 'SIMD within a register' or the 'bit-slicing' technique. The first mention of bit-slicing technique applied to cryptographical problem belongs, seemingly, to E.Biham [?].

Now we describe the main idea of the bit-slicing technique. Consider an arbitrary total boolean function

$$f : \{0, 1\}^n \rightarrow \{0, 1\}$$

This function can be represented in the form of the Boolean circuit $C(f)$ over some complete basis B . A common example of such basis is $B = \{\wedge, \vee, \neg\}$, but we'll use the basis $B = \{\wedge, \vee, \neg, \oplus\}$ instead, since it better fits our goals.

Now consider the problem of calculating the arbitrary total boolean function $f : \{0, 1\}^n \rightarrow \{0, 1\}$ over all 2^n possible inputs. For each input $X \in \{0, 1\}^n$ one can calculate the value of f as a superposition of the basis functions, according to the scheme $C(f)$. We can now select at least one fixed order of calculation of basis functions from $C(f)$, that results in getting the value of f . Let m be the number of internal nodes in $C(f)$. Assuming that the calculation of one basis function takes one processor instruction, the computation of f over all inputs from $\{0, 1\}^n$ will take $m \cdot 2^n$ instructions.

A computational architecture that, in a single instruction execution, calculates many copies of the same function over many different memory cells, is called the 'single instruction, multiple data' (SIMD) architecture. When a modern computational device executes a bitwise logical instruction over its GPRs, it effectively acts as a SIMD device, individual bits of GPRs playing the role of the individual memory cells of a SIMD device. The calculation order of functions in the scheme $C(f)$ always stays the same. This makes it possible to compute this functions over as many inputs, as there are bits in the device's GPR. In effect, if D is the device's GPR capacity, we can simultaneously walk D instances of the scheme $C(f)$, effectively calculating the value of f for inputs X_1, \dots, X_D .

Let's examine the arbitrary basis function g with arity 2, and the corresponding internal node of the scheme $C(f)$. We denote it as $G(x_1, x_2)$, meaning that it has a single output and two inputs, which values are determined by boolean variables x_1, x_2 . Next, collate g with three GPRs denoted $R_1(g), R_2(g), R_3(g)$, each of which is comprised of D single-bit memory cells, filled in the following way:

- Register R_1 contains D values of the variable x_1 , corresponding to X_1, \dots, X_D ;
- Register R_2 contains D values of the variable x_2 , corresponding to X_1, \dots, X_D ;

- Register R_3 contains D values of the function g , corresponding to matching values of X_1, \dots, X_D .

Suppose that all D instances of values of g (matching the corresponding values of x_1 and x_2) can be computed as a result of a single bitwise instruction applied to register $R_1(g), R_2(g)$, while their result is put into register $R_3(g)$. If this fact is holding true for every basis function in the scheme $C(F)$, the computation of f for every input from $\{0, 1\}^n$ will require $m \cdot \frac{2^n}{D}$ instructions. This is the key idea of the bit-slicing technique [?].




Fig. 3. Simultaneous computation of D instances of scheme $C(f)$ with bit-slicing technique

We will call the process of computation of the function f on a single input from $\{0, 1\}^n$ a "thread", by analogy with the computational threads in a SIMD device. So, with the use of bit-slicing technique, it takes m instructions to complete the computation of D threads.

Now we describe the details of bit-slicing implementation of the A5/1 generator. Suppose that a computational device is able to calculate D instances of any

function from the basis $B = \{\wedge, \vee, \neg, \oplus\}$. Each of $n \in \{1, \dots, 64\}$ cells of LFSR 1-3 gets a matching word $W_n \in \{0, 1\}^D$:

$$LFSR1 : W_1, \dots, W_{19};$$

$$LFSR2 : W_{20}, \dots, W_{42};$$

$$LFSR3 : W_{43}, \dots, W_{64}.$$

In a bit-slicing technique, the shift of the LFSR register (LFSR 1 in this example) will take the following form:

$$W'_1 = W_{19} \oplus W_{18} \oplus W_{17} \oplus W_{14},$$

$$W'_n = W_{n-1}, n \in \{2, \dots, 19\},$$

where \oplus is the bitwise addition modulo 2 of boolean vectors of length D operator. The calculation of the keystream bit will look like:

$$W_{out} = W_{19} \oplus W_{41} \oplus W_{64}.$$

The conditional clocking is somewhat more complex to implement in a bit-slicing technique. First, to know if the LFSRs should be shifted or not, one needs to calculate the corresponding clocking flags F_1, F_2, F_3 using the majority function:

$$W_{maj} = maj(W_8, W_{30}, W_{52}) = (W_8 \wedge W_{30}) \vee (W_8 \wedge W_{52}) \vee (W_{30} \vee W_{52}),$$

$$F_1 = W_8 \oplus \neg W_{maj},$$

$$F_2 = W_{30} \oplus \neg W_{maj},$$

$$F_3 = W_{52} \oplus \neg W_{maj},$$

(all operations are bitwise operations on the vector of the length D).

To implement the conditional clocking of a LFSR one can use the bitwise counterpart of the "bitselect" function of arity 3:

$$a, b, c \in \{0, 1\};$$

$$BS(a, b, c) = \begin{cases} b, & a = 1, \\ c, & a = 0. \end{cases}$$

If the computational architecture lacks the hardware implementation of this function, it can be emulated with the usage of the standard bitwise functions corresponding to the functions of the basis B (example for LFSR 1):

$$BS(a, b, c) = (a \wedge b) \vee (\neg a \wedge c).$$

Clocking the LFSR 1 with the use of $BS(a, b, c)$ and corresponding flag F_1 looks the following way:

$$W'_1 = BS(F_1, (W_{19} \oplus W_{18} \oplus W_{17} \oplus W_{14}), W_1);$$

$$W'_n = BS(F_1, W_{n-1}, W_n), n \in \{2, \dots, 19\}.$$

Some important details about the bit-slicing implementation of the Anderson's attack should be covered still. The Anderson's attack follows 2 stages:

1. Calculating the values of 11 bits of LFSR 2 lying left of the clocking bit with the information from the remaining 53 bits (guessed) and the keystream (known).
2. Clocking the generator as normal to check if the guessed filling of the generator matches the known keystream.

The irregular clocking of the A5/1 generator makes it generally impossible to predict how many clockings of the whole generator (bits of keystream) would be needed to shift LFSR 2 11 times to complete the Stage 1 of the attack. Therefore, we again put to use the bitselect function to implement the split of the attack into 2 stages. Since each individual thread should be able to stay in a stage and advance to the next stage independent of other threads, we introduce the special boolean vector $\phi = (\phi_1, \dots, \phi_D)$, called the "attack stage flag". The thread with number $i, i \in \{1, \dots, D\}$ being in the Stage 1 of the attack corresponds to $\phi_i = 0$, and the Stage 2 of the attack corresponds to $\phi_i = 1$. Let y_1, \dots, y_{64} be the bits of the keystream analyzed. Now the clocking of the LFSR 2 considers the stage of the attack through the usage of the attack stage flag:

$$\begin{aligned} W_{41}^* &= BS(\phi, W_{41}, (y \oplus W_{19} \oplus W_{64})); \\ W_{20}' &= BS(F_2, (W_{41}^* \oplus W_{40}), W_{20}); \\ W_n' &= BS(F_2, W_{n-1}, W_n), n \in \{20, \dots, 41\}. \end{aligned}$$

Here W_{41}^* is a helper vector holding temporary data. y is the current bit of the keystream, in the form of a vector consisting of D copies of k -th bit of the keystream. In effect, Stage 1 of the attack's goal is to calculate the last bit of the LFSR 2 from known keystream bits and known (guessed) last bits of the LFSRs 1 and 3. At the Stage 2 the last bit of the LFSR 2 is known and the whole generator scheme is clocked as normal. For the i -th thread the attack stage flag ϕ_i is set to 1 after the thread's instance of LFSR 2 was shifted 11 times. To count the number of LFSR 2 shifts on a per-thread basis, the bit-slicing implementation of an incremental counter is used.

The Anderson's attack algorithm described above was realized on an NVIDIA GPU with the use of CUDA SDK 8.0 [?]. The comparison of the performance of a GPU to a CPU in execution of bit-slicing and LFSR precomputation-based [?] implementations of Anderson's attack is shown in [?].

Data provided in the Table 1 tells us that even one mid-range consumer GPU is enough to make the Anderson's attack run time practical (it will take around 250 hours). A modern computational cluster outfitted with GPUs will complete the attack in mere minutes or less. The Anderson's attack's advantage over the rainbow-tables attack [?] is the former's ability to restore the secret key from 64 bits of keystream with the 100% probability. It's advantage over the attack described in [?] is in a usage of a consumer-grade off-the-shelf hardware.

Table 1. (LOP3.LUT is a special instruction that implements arbitrary bitwise arity 3 functions in hardware. We use it as a substitute for the "bitselect" instruction.)

Performance of the different implementations of the Anderson's attack (search space is 2^{53}) on a CPU and a GPU, measured in millions of key checks per second.

Computational device	Bit-slicing	LFSR precomputation [?]
CPU Intel Core I7 930	37	7
GPU NVIDIA GTX 1050 Ti	9180	483
GPU NVIDIA GTX 1050 Ti (LOP3.LUT)	11950	-

4 Implementation of the Anderson's Attack in a volunteer computing project

In order to solve 10 cryptanalysis problems for the A5/1 keystream generator we launched the volunteer computing project AndersonAttack@home. The computing application of this project is based on the CUDA implementation, which was described in the previous section.

Volunteer computing is a type of distributed computing which uses computational resources of PCs (hosts) of private persons called volunteers. Each volunteer computing project is designed to solve one or several hard problems. When PC is connected to the project, all the calculations are performed automatically and do not inconvenience user since only idle resources of PC are used. Volunteers get credits (points) for each calculated task.

Volunteer computing project consists of the following basic parts: server daemons, database, web site and client application. Daemons include WORK GENERATOR (generates tasks to be processed), VALIDATOR (checks the correctness of the results received from volunteer's PCs) and ASSIMILATOR (processes correct results). Client application should have versions for the widespread computing platforms. One of attractive features of volunteer computing is its low cost — to maintain a project one only needs a dedicated server working 24/7.

AndersonAttack@home is based on BOINC (Berkley Open Infrastructure for Network Computing [?]), which is the most popular platform for volunteer computing. Each BOINC-based project is separate, but any private person can take part in arbitrary project with the help of standard BOINC manager. This BOINC manager allows users to change projects priority and to regulate the mode of PC using.

BOINC provides a form of redundant computing. According to this, several (at least 2) similar tasks are created for each workunit. These tasks are processed on multiple hosts, the results are compared on a project server by validator, and are accepted only when a 'consensus' is reached. The sufficient number of successful results is called a 'quorum'. If for a particular workunit results of tasks are nonsimilar, then new tasks must be created and sent. The scheme of a BOINC-based project with the quorum of 2 is shown on the Fig. 1.

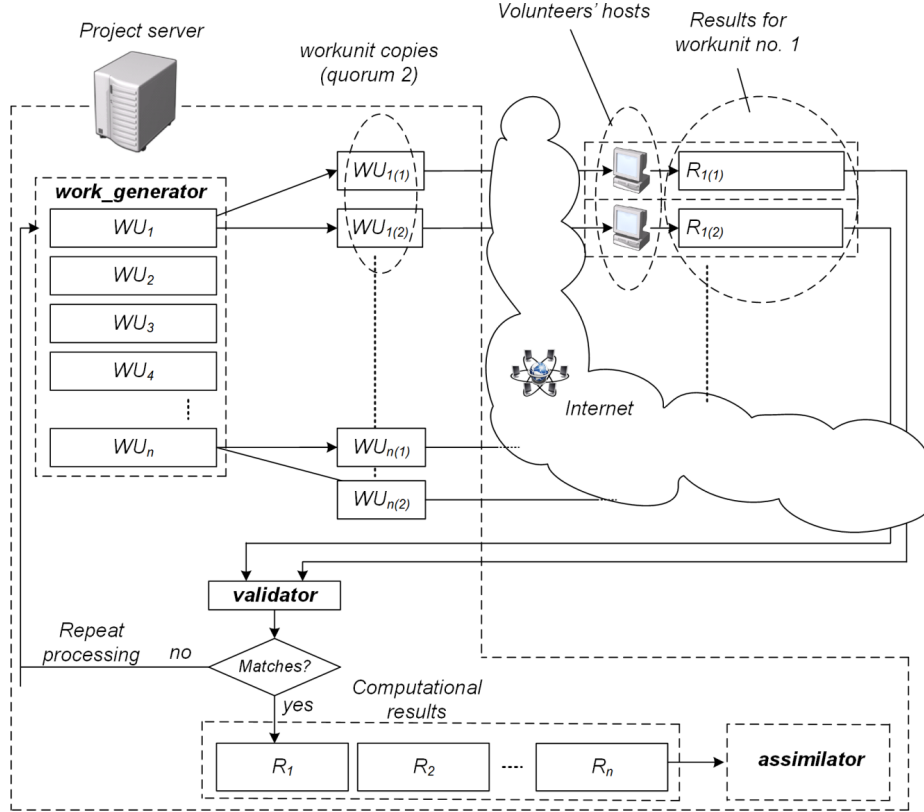


Fig. 4. Simultaneous computation of D instances of scheme $C(f)$ with bit-slicing technique

On the first stage of our experiment on AndersonAttack@home, a family of workunits was generated on the project server. In each workunit values of 12 variable cells out of 53 (see the previous section) were fixed. Thus, 40960 workunits were generated for 10 cryptanalysis instance in total. Usually the value of deadline for workunits in BOINC projects is equal to 10-14 days. In our project we used the deadline of 1 day, because the experiment was quite small. On the next stage all generated workunits were processed in a desktop grid formed by project's hosts. This took about 7 days. As a result, solutions for all considered instances were successfully found. All these solutions are listed in Table 1. It should be noted, that collisions were found for 7 out of 10 instances (5 of them have 1 collision, each of other 2 instances has 2 collisions).

Table 2.

Original secret keys and collisions of the generator A5/1 (in hexadecimal format) which were found in AndersonAttack@home.

Instance	Keystream	Secret key
1	0x770c0410869366f1	original 0x11b8e4340276c4ee collision1 0x42634f3266d302a3
2	0xae9590560c26e9ed	original 0x4c656fd73e59ab9b collision1 0xcf23e4722e3cfb68
3	0xdd4b3ab7f6cf8224	original 0x09429d158555f4b3 collision1 0x09429d158553e967 collision2 0x40e5f2c8128a1781
4	0x93cd42d97eb75fd9	original 0xfa386a338355aafd collision1 0xf9e81096bb4d0aad collision2 0xf9e81096bb4a8556
5	0x925e423c98121152	original 0xe5cf81035ce5fbe2
6	0x3b3464bd6e377b87	original 0x9625e9d810b46248 collision1 0xf5aalbe2d6c36e18
7	0x0367d29121dd1677	original 0xd1b8b06086edf162
8	0x6b49230b7fc0249d	original 0xbe81a896968c486b
9	0xc65847556752d14c	original 0xb6f65d2855a211c0 collision1 0xb6f65d2855a508e0
10	0x07bb7f83d26072ec	original 0x122a1a2955286b9f collision1 0xd5151aaa50490012

According to BOINC statistics, 143 active hosts of 90 volunteers participated in the experiment. Here by "active" we mean a host which correctly processed at least one task. Active volunteer has at least one active host. It should be noted, there is only 1 volunteer from Russia in top-10 project volunteers list (formed according to the amount of credits), the most of them are from United States (4 persons). In top-10 list of teams there is only one team from Russia too.

5 Related works.

As was noted before, at the time of writing of our article the A5/1 algorithm remains one of the most popular research subjects among cryptography specialists, along with DES and RSA algorithms. The exact authorship of this algorithm is unknown, but some experts argue that chances are high that its creation should be attributed to a French special agencies. Numerous sources describe the events that led to its structure leaked to general public. In fact, its internal workings were already known to research community in 1994: R.Anderson's note (a mailing list message) written in that period was one of the first works on A5/1 cryptanalysis. The complete knowledge of the A5/1 were acquired in 1999 as a result of a reverse engineering of a mobile phone.

Earlier we mentioned Anderson's attack (described by Ron Anderson in a short note in 1994 [?]) was the first attack on A5/1 with complexity less than that of a trivial brute-force walk over the whole key space. Shortly afterwards J.Golic suggested the "Alleged" attack on the A5/1, based on a linearization of the equations describing A5/1 [?]. The Golic's attack estimated complexity is C^2 ⁴⁰, where C denoting the complexity of solving the system of linear equations over $GF(2)$ with a very substantial number of dimensions. The next widely known attack on the A5/1 was presented in [?]. To speed up the A5/1 algorithm its authors used the technique of precomputation of LFSRs we mentioned in [?]. It's worth to note that the attack presented in [?], along with those presented in [?], [?], [?], require substantial lengths of known keystream (several seconds at best). These attacks can be deemed realistic, but they do not demonstrate vulnerability of the GSM security protocols as evidently, as those belonging to the "advanced brute force" class, that we will discuss below. For the rest of the section we consider an attack practical if its real-life implementation is able to persistently find the secret key by analyzing no more than 2 frames (912 bits) of the keystream, the amount that could always be extracted from a call in a GSM network by exploiting the protocol vulnerability described in Section 2.

It seems that the first practical attack ('practical' in a sense we established above) was presented in [?]. Its authors implemented the optimized variant of the Anderson's attack on specialized computational device of their own design, assembled from 120 "Xilinx Spartan 3" FPGAs. The authors state [?] that the attack took about 6 hours to complete. It's worth mentioning that the COPACOBANA architecture was used for cryptanalytical research of several other ciphers, e.g. DES [?].

The first estimates of the time required for the cryptanalysis of A5/1 on a computational cluster in the form of a boolean satisfiability problem (SAT) [?] were presented in [?]. SAT-based cryptanalysis is considered a perspective direction of cryptanalysis that seems to be on the rise for the last decade. It operates within a paradigm that the cryptanalytical problem could be efficiently represented in the form of a SAT problem. This could be achieved with the help of special translators such as [?], [?], [?]. After the problem was translated to a SAT form, it can be efficiently solved by modern SAT algorithms that could run in a distributed computing environment. The cryptanalysis time of

A5/1 estimated in [?] was realistic, but since the exclusive usage of the whole computational cluster for prolonged period of time was not an option, the actual attack was performed a year later in a specialized grid-system [?]. In the following years these results were improved - the search space of 2^{31} SAT instances was processed completely and collisions (different secret keys generating the same keystream) have been found [?]. At the same period it became apparent that the organizational requirements of this kind of cryptanalytical attack fits well into the ideology of volunteering computing projects. To bring this line of thought to life in 2012 SAT@home [?] volunteering computing project was inceptioned. Through 4 years of it's activity a number of hard combinatorial problems were solved within the SAT paradigm. Among those were several dozens of cryptanalysis of A5/1 generator problems (only one burst - 114 bits of keystream - was used each time). These results were published in [?], [?].

The work [?] included the estimates of creation time and disk space that would be required to create the rainbow-tables that would make possible to 'break' the A5/1 on an average PC in several minutes. [?] estimated these to occupy about 7 Tb of disk space. "The A51 Cracking Project" put the significantly more compact (2 Tb) tables into public domain at the end of 2009. By analysing 2 frames (912 bits) of known keystream with the help of this tables one can restore the secret key with probability of success over 85%. This method of cryptanalysis of the A5/1 algorithm could be assumed to be the most practical. The fact that it does not provide 100% guarantee of success from analysis of 2 frames does not matter in real circumstances, where much longer samples of keystream are available to the attacker. It should be noted that if the tables are lost, it will require a vast amount of computational resources to generate them again.

Finally, the attack we presented in this paper does not lose on the practical side to the rainbow-tables attack, because it allows to complete the cryptanalysis with 100% success rate in a period of several hours in a small volunteering computation project. Moreover, the attack requires only 64 bits of known keystream. Volunteering computing projects are rising in popularity and publicity in the last years, so any "hot" problem (like a cryptanalysis of one of the most commonly used ciphers) immediately gets volunteers attention.

As a conclusion, we would like to express the hope that the efficiency of our implementation of attack on A5/1 generator will serve as yet another strong argument against the usage of this algorithm in the transfer of sensible data through GSM networks.

References